




Czech Technical University in Prague
Faculty of Nuclear Sciences and Physical
Engineering



Distributed Kalman filtration under unknown spatially heterogeneous noise

Master Thesis

Author:	Bc. Daniel Hnyk
Supervisor:	Ing. Kamil Dedecius, Ph.D.
Academic	year: 2017/2018

Acknowledgment:

add ack

Declaration:

I hereby certify that this text represents my own work and that all used sources and materials are listed in the bibliography.

Prague, Friday 20th April, 2018

Daniel Hnyk

Title: Distributed Kalman filtration under unknown spatially heterogeneous noise

Author: Bc. Daniel Hnyk

Branch of study: Mathematical Informatics

Abstract:

add abstract

Key words: Bayesian inference, Kalman filter

Contents

Introduction	6
1 Bayesian probability	7
1.1 Probability interpretations	7
1.2 Why Bayesianism	7
1.3 Principles of Bayesian probability	8
1.4 Conjugate distributions and exponential family	8
1.5 Statistical model	8
1.6 Statistical inference	9
1.7 Variational Bayes inference	9
1.8 Coordinate ascent mean-field variational	11
1.9 Example	12
2 Distributed Kalman filter	14
2.1 Introduction	14
2.2 Description	15
2.3 Mathematical definition	15
2.4 Note on optimality	16
2.5 Noise analysis	17
2.6 Methods	17
2.7 Distributed Kalman Filter	17
3 Distributed Variational Bayes Adaptive Kalman Filtration	18
3.1 Problem formulation	18
3.2 Variational Bayes	21
3.3 Algorithm summary	24
4 Results	26
4.1 Implementation	26
4.2 Simulations	27
Conclusion	42
Bibliography	42

Introduction

Add introduction

Chapter 1

Bayesian probability

In this chapter, I summarize basics and preliminaries of Bayesian interpretation of probability and shortly compare it to frequentist approach. Reader is expected to know basic concepts from probability such as probability event, space, measure, product and sum rule, conditional and joint probability. All of these can be found in [1].

1.1 Probability interpretations

As it is a case with other branches of mathematics, probability theory has multiple interpretations with their own strengths and weaknesses. In this particular case, it is a subject of dispute if the term “interpretations” should be used, since there is no single formal system “probability”. For example, some of the leading interpretations fail to satisfy the most common Kolmogorov’s axiomatization and they still prove themselves useful[2]. Other axiomatization has been proposed for different interpretation [3] and some of the leading interpretations such as [1] don’t even derive from *axioms*, but rather *desiderata*¹. Interpretations don’t differ only in the underlying set of axioms, but also in their practical usage or their epistemological status. Mathematicians as well as philosophers has been tackling these issues at least for over five centuries.

1.2 Why Bayesianism

It’s hard to fully appreciate and understand main concepts and advantages of Bayesian reasoning if one is not familiar with other interpretations, such as frequentist. Its full description can be found in [4], here we’ll cover just the basics arising from the comparison.

The main idea of Bayesian interpretation is its subjective nature, where there is no such thing as *real* probability in the world, but probabilities are in actor’s minds. Probability then represents agent’s *state of knowledge* or *degree of belief*[5, 6, 7]. This is not a negligible formal difference – it has important consequences, such as the fact that all agents with different state of knowledge can assign different probabilities to the same event. In Bayesian framework, one explicitly convey uncertainty in his models and objects (such as random variables).

Another important aspect worth mentioning is also the fact that Bayes interpretation allows doing predictions for events which haven’t occurred yet (where frequentists are unable to do anything). This property is also connected to a fact that one can do partial (*online*) updates based on new evidence.

¹(I) representations of plausibility are to be given by real numbers; (II) plausibilities are to be in qualitative agreement with common sense; and (III) the plausibilities are to be “consistent”, in the sense that anyone with the same information would assign the same real numbers to the plausibilities

1.3 Principles of Bayesian probability

Core idea of Bayesian interpretation is surely a famous Bayes theorem 1.1, which is easily obtained from definition of conditional probability [8].

Theorem 1. Let P be a probability measure, \mathbf{A} and \mathbf{B} are events, $p(\mathbf{B}) \neq 0$. Then

$$p(\mathbf{A}|\mathbf{B}) = p(\mathbf{A}) \cdot \frac{p(\mathbf{B}|\mathbf{A})}{p(\mathbf{B})} \quad (1.1)$$

where:

1. $p(\mathbf{A})$ and $p(\mathbf{B})$ are the probabilities of observing \mathbf{A} and \mathbf{B} independently.
2. $p(\mathbf{A}|\mathbf{B})$ is the probability of observing event \mathbf{A} given that \mathbf{B} is true.
3. $p(\mathbf{B}|\mathbf{A})$ is the probability of observing event \mathbf{B} given that \mathbf{A} is true. This factor is called **likelihood**.

To relate this simple formula to previous section, consider a scenario where one wants to use Bayesian framework for *inference*. Suppose that \mathbf{A} is a hypothesis whose probability can be affected by some evidence and \mathbf{B} is an evidence for this hypothesis. $p(\mathbf{A})$ in formula 1.1 can then be seen as a **prior probability** – a probability one would assign to \mathbf{A} is true *before* observing the evidence. Factor $p(\mathbf{B}|\mathbf{A})$, the *likelihood*, represents how likely we expect to observe \mathbf{B} *given* that \mathbf{A} is true. $p(\mathbf{B})$ is same for all possible hypothesis and serves as a normalizing factor (that is the reason why is often left out in many computations). And finally, $p(\mathbf{A}|\mathbf{B})$, which is the final **posterior probability** which is the factor one usually care about – the probability of \mathbf{A} after observing \mathbf{B} .

A factor $\frac{p(\mathbf{B}|\mathbf{A})}{p(\mathbf{B})}$ has also quite useful interpretation and that is the overall impact of evidence \mathbf{B} on the probability of \mathbf{A} .

1.4 Conjugate distributions and exponential family

Borrow this part from Mr Dedecius article page 3 - source code and cite it. There is informal definition right before Definition 2 and then consecutive much more detailed definitions. Also use rationale from before definition 2.

The updates can be significantly easier if *conjugate distributions* are used instead.

1.5 Statistical model

Chosen interpretation has a nontrivial influence on what such concept as *statistical model* is and how *inference* is done. I will follow traditional approach taken in [9, 10] and extend it of some necessary concepts from [11].

A statistical model is a pair of (S, M) , where S is a set of possible observations (*sample space*) and M is a set of probability distributions on S . Distribution in M are expected to be approximately close to the “true” distribution which generates the data². One of the main distinctions which can be made to separate families of statistical models is how distributions in M are described:

²There is a saying: “All models are wrong, but some of them are useful”

1. **Parametric models:** in this case, the set M is parameterized by some parameter θ and can have values from *finite* parameter space $\Theta \subseteq \mathbb{R}^d$. In other words: $M = \{P_\theta | \theta \in \Theta\}$.
2. **Nonparametric models:** this family is an unfortunate misnomer – nonparametric models also have parameters, but they differs from parametric models in a way that the parameter space is Θ is *infinite* – it is not fixed and can grow with the amount of data.

In Bayesian setting, parameters of the model have some prior distributions assigned. These distributions capture prior knowledge one posses before the application. For instance, this can be based on previous research or reliable enough observations. If there is no prior knowledge available, so called *uninformative* priors are being used instead. *Uninformative* is again a misnomer – what is meant by these is the fact that they are not subjectively elicited. Fully Bayesian model is then a model in which *all* parameters have some prior assigned. An interesting property of Bayesian approach arise – random variables (including latent variables) and parameters are treated in the same way.

1.6 Statistical inference

Process of inferring properties of the data underlying distribution is called statistical inference[12]. Formally, it's justification of restricting parameter space based on the data (by e.g. choosing a point estimate for a given parameter).

Unfortunately, for many real world scenarios, the approach of directly inferring the properties is not possible. One of many reasons can be to high dimensionality of the task or the form of the posterior distribution may be to complex, let alone the fact that it does not have to have analytical solution[13]. In these scenarios, approximation methods are being used instead.

1.6.1 Approximation methods

Using a breakdown from [13], approximation methods can be divided to two groups based on their stochastic or deterministic behavior. A notable example of the former is Markov chain Monte Carlo[13, (Chap.11)]. A representative approach of the latter is e.g. variational inference described in detail in section 1.7. The main difference is the fact that stochastic schemes, *given infinite resources*, are guaranteed to get the exact results and the approximation arise from the finite amount of those resources. On contrary, deterministic approximation schemes are based on *analytical approximations* to the posterior distribution and hence are generally not able to generate exact results. Both approaches are then complementary to each other, one being useful for situation where the second is unsuitable and vice versa.

1.7 Variational Bayes inference

In this work, a family of deterministic approximation methods *variational inference* (or *variational Bayes*) is used, excellently described in [11, (p. 463)] and [14]. This technique is based on using solution of an optimization problem to statistical inference, more exactly finding an input which minimize a specific functional. As was described in 1.6.1, this method is guaranteed to get an exact result given some family of possible input functions over which one minimizes. The approximation arise from limiting the possible inputs, for instance by considering only quadratic functions or, as is widely used and is also our case, functions which factorizes in a specific way.

Let $p(\mathbf{X}, \mathbf{Z})$ be a fully Bayesian model, where all prior distributions are given. Let denote the set of all parameters and all latent $\mathbf{Z} = \{F_1, \dots, F_N\}$ and the set of all observed variables

$\mathbf{X} = \{\curvearrowright_K, \dots, \curvearrowright_N\}$. The goal of the inference is to find the posterior distribution $p(\mathbf{Z}|\mathbf{X})$ and the distribution of the model evidence $p(\mathbf{X})$ ³. Unfortunately, in many real world scenarios, $p(\mathbf{Z}|\mathbf{X})$ is almost always intractable (by e.g. trying to integrate all configurations of the hidden variables in denominator). That is where *variational lower bound* comes in. Instead of trying to compute $p(\mathbf{Z}|\mathbf{X})$ directly, we consider $q(\mathbf{Z})$ which is as close approximation as possible to the former and has a convenient and tractable form (their expectations are computable). Furthermore, these approximate distributions can also have their own *variational parameters* which are considered to be in \mathbf{Z} as well. As a measure of closeness between the approximate $q(\mathbf{Z})$ and $p(\mathbf{Z}|\mathbf{X})$, Kullback-Leibler (KL) divergence 1.2 is used⁴.

$$\text{KL}(q||p) = - \int q(\mathbf{Z}) \ln \frac{p(\mathbf{Z}|\mathbf{X})}{q(\mathbf{Z})} d\mathbf{Z} \quad (1.2)$$

In a case when $p(\mathbf{X}|\mathbf{Z}) = q(\mathbf{Z})$, Kullback-Leibler divergence is effectively zero. The goal is the following optimization problem:

$$q(\mathbf{Z}) = \arg \min_{q(\mathbf{Z})} \text{KL}(q(\mathbf{Z})||p(\mathbf{Z}|\mathbf{X})). \quad (1.3)$$

The optimization 1.3 cannot be performed directly. To compute $\text{KL}(q||p)$, the unknown evidence $p(\mathbf{X})$ is needed, as can be seen from the following derivation

$$\text{KL}(q||p(\mathbf{Z}|\mathbf{X})) = \mathbb{E}[\ln q] - \mathbb{E}[\ln p(\mathbf{Z}|\mathbf{X})] = \mathbb{E}[\ln q] - \mathbb{E}[\ln p(\mathbf{Z}, \mathbf{X})] + \ln p(\mathbf{X})$$

Instead of trying to compute $\text{KL}(q||p)$, we are going to optimize an alternative objective that is equivalent to $\text{KL}(q||p)$ up to an added constant. Let us then define *variational lower bound* (also called *ELBO*) as follows:

$$\mathcal{L}(q) = \int q(\mathbf{Z}) \ln \frac{p(\mathbf{X}, \mathbf{Z})}{q(\mathbf{Z})} d\mathbf{Z} = \mathbb{E}[\ln p(\mathbf{Z}, \mathbf{X})] - \mathbb{E}[\ln q(\mathbf{Z})]. \quad (1.4)$$

It can be shown that $\mathcal{L}(q)$ has an important property being a lower bound of the log probability $\ln p(\mathbf{X}) \geq \mathcal{L}(q)$. Henceforth, when one wishes to maximize marginal $p(\mathbf{X})$, he can instead maximize its variational lower bound $\mathcal{L}(q)$.

Combining 1.4 and 1.2, the following relationship can be derived:

$$\ln p(\mathbf{X}) = \mathcal{L}(q) + \text{KL}(q||p). \quad (1.5)$$

Since $p(\mathbf{X})$ doesn't depend on q , maximizing $\mathcal{L}(q)$ is possible and it is equivalent to minimizing $\text{KL}(q||p)$.

Again, since we don't know the true posterior distribution, we work with some family of distributions $q(\mathbf{Z})$ for which $\mathcal{L}(q)$ becomes tractable and search for the candidate which maximizes it. Obviously, the choice of the distribution family is critical here – while it needs to be tractable, it still needs to be flexible enough (as much as possible) to provide *accurate enough* approximation, goals usually going against each other. This is the part from where approximation arise in variational Bayes as has been discussed in 1.6.1.

³ $p(\mathbf{X}) = \int p(\mathbf{Z}, \mathbf{X}) d\mathbf{Z}$, which is usually unavailable

⁴We omit arguments of distributions where possible for readability. Most often: $q := q(\mathbf{Z})$ and $p = p(\mathbf{Z}|\mathbf{X})$

1.7.1 Factorized distributions

A common choice for such a family of distributions described in the chapter 1.7 has been *factorized distributions* formalized in physics theory called *mean field theory* [15]. The assumption on the family is quite simple – we treat each variable in \mathbf{Z} as independent, in other words, the family **factorizes**:

$$q(\mathbf{Z}) = \prod_{i=1}^M q_i(\mathbf{Z})$$

It is clear that the assumption is often not met. Variables are dependent in the real world – it is the reason why it is so hard to obtain the posterior distribution directly.

Notice that we do not require any specific form of $q_i(\mathbf{Z}_i)$. Also, it is worth noting that this is *not* a model of the observed data – it is the ELBO and KL minimization problem, which connects this to the model and data.

1.8 Coordinate ascent mean-field variational

Coordinate ascent mean-field variational (CAVI) is a commonly used algorithm for solving optimization problem 1.3[14][13].

1.8.1 Derivation

Optimal $q_i^*(\mathbf{Z}_i)$

First, consider the complete conditional of \mathbf{Z}_i , which is $p(\mathbf{Z}_i|\mathbf{Z}_{-i}, \mathbf{X})$. If we fix all other variational factors $q_l(\mathbf{Z}_l), l \neq i$, the optimal $q_i(\mathbf{Z}_i)$ satisfies the following⁵

$$q_i^*(\mathbf{Z}_i) \propto \exp\{\mathbb{E}_{-i}[\ln p(\mathbf{Z}_i|\mathbf{Z}_{-i}, \mathbf{X})]\} \propto \exp\{\mathbb{E}_{-i}[\ln p(\mathbf{Z}_i, \mathbf{Z}_{-i}, \mathbf{X})]\} \quad (1.6)$$

where in the last term, we use the fact that $q_{-i}(\mathbf{Z}_{-i}) = \prod_{l \neq i} q_l(\mathbf{Z}_l)$.

Obtaining KL

Let us now rewrite $\mathcal{L}(q)$ in 1.4 as a function of $q_i(\mathbf{Z}_i)$

$$\mathcal{L}(q_i) = \mathbb{E}_i[\mathbb{E}_{-i}[\ln p(\mathbf{Z}_i, \mathbf{Z}_{-i}, \mathbf{X})]] - \mathbb{E}_i[\ln q_i(\mathbf{Z}_i)] + C \quad (1.7)$$

where in the first term is just an iterated expectation and in the second we retain only product term containing $q_i(\mathbf{Z}_i)$, while other parts can be moved to a constant C thanks to using factorized distribution.

Now it's easy to see that 1.7 is a negative KL divergence between $q_i(\mathbf{Z}_i)$ and $q_i^*(\mathbf{Z}_i)$. That gives as all necessary part for the CAVI algorithm as described in [14]:

Algorithm 2. CAVI

Input: A model $p(\mathbf{X}, \mathbf{Z})$, a dataset \mathbf{X}
Output: A variational density $q(\mathbf{Z}) = \prod_{i=1}^m q_i(\mathbf{Z}_i)$
Initialize: Variational factors $q_i(\mathbf{Z}_i)$
while the $\mathcal{L}(q)$ has not converged
 for $i \in \{1, \dots, m\}$ **do**
 Set $q_i(\mathbf{Z}_i) \propto \exp\{\mathbb{E}_{-i}[\ln p(\mathbf{Z}_i, \mathbf{Z}_{-i}, \mathbf{X})]\}$

⁵By $\mathbb{E}_{-i}[\dots]$, we denote expectation over all q distributions *except* q_i . Similarly, \mathbf{Z}_{-i} means all latent variables except \mathbf{Z}_i .

```

end
Compute  $\mathcal{L}(q) = \mathbb{E}[\ln p(\mathbf{Z}, \mathbf{X})] - \mathbb{E}[\ln q(\mathbf{Z})]$ 
end
return  $q(\mathbf{Z})$ 

```

It is guaranteed that this algorithm converge to a local minimum.

1.9 Example

Combination of all of the above in a single example. Ideally with nice visualization.

Let us demonstrate the above with an example from [13].

The model Consider the i.i.d dataset $\mathbf{X} = \{x_1, \dots, x_N\}$ generated by (unknown) Gaussian distribution $\mathbf{X} \sim \mathcal{N}(\mu, \tau^{-1})$. The goal is to infer posterior $p(\mu, \tau | \mathbf{X})$ and the joint probability is

$$p(\mathbf{X}, \mu, \tau) = p(\mathbf{X} | \mu, \tau) p(\mu | \tau) p(\tau)$$

Factorized distribution Let us now approximate $p(\mu, \tau | \mathbf{X})$ by $q(\mu, \tau)$ with the assumption that q factorizes: $q(\mu, \tau) = q(\mu)q(\tau)$.

Conjugated priors The complete data likelihood is Gaussian distribution and hence in exponential family. We set a non-informative conjugated priors for the μ and τ as follows

Is it correct? Is it likelihood to which we set conjugated priors?

$$\begin{aligned}\mu &\sim \mathcal{N}(\mu_0, (\lambda_0 \tau)^{-1}) \\ \tau &\sim \text{Gamma}(a_0, b_0)\end{aligned}$$

where the hyperparameters $\mu_0, \lambda_0, a_0, b_0$ are initially set to some small positive number.

Form of $q(\mu)$ and $q(\tau)$ This is usually the hardest part of variational inference. The derivation of the following terms is according to 1.6 and in detail can be found in [13].

The derivations yields following for $q_\mu^*(\mu)$

$$\begin{aligned}q_\mu^*(\mu) &\sim \mathcal{N}(\mu | \mu_N, \lambda_N^{-1}) \\ \mu_N &= \frac{\lambda_0 \mu_0 + N \bar{x}}{\lambda_0 + N} \\ \lambda_N &= (\lambda_0 + N) \mathbb{E}_\tau[\tau] = (\lambda_0 + N) \frac{a_N}{b_N} \\ \bar{x} &= \frac{1}{N} \sum_{n=1}^N x_n\end{aligned}$$

and for $q_\tau^*(\tau)$:

$$\begin{aligned}
q_\tau^*(\tau) &\sim \text{Gamma}(\tau \mid a_N, b_N) \\
a_N &= a_0 + \frac{N+1}{2} \\
b_N &= b_0 + \frac{1}{2} \mathbb{E}_\mu \left[\sum_{n=1}^N (x_n - \mu)^2 + \lambda_0 (\mu - \mu_0)^2 \right] = \\
&= b_0 + \frac{1}{2} \left[(\lambda_0 + N) (\lambda_N^{-1} + \mu_N^2) - 2 \left(\lambda_0 \mu_0 + \sum_{n=1}^N x_n \right) \mu_N + \left(\sum_{n=1}^N x_n^2 \right) + \lambda_0 \mu_0^2 \right]
\end{aligned}$$

Perform updates based on CAVI Initially, compute $N, \sum_{n=1}^N x_n, \sum_{n=1}^N x_n^2$ and set $\lambda_{N,t}$ to some random value. Then iterate these steps until convergence:

1. plug in $\lambda_{N,t}$ to obtain $b_{N,t}$
2. compute new $\lambda_{N,t+1}$ based on $b_{N,t}$ and all other terms needed

Is it correct that a_N, μ_N doesn't update? The "N" suffix here seems to be quite redundant.

Results After m iterations, we are now having new values of all parameters including hyper-parameters. Hence we are able to approximate the posterior distribution. The joint probability is

$$\begin{aligned}
p(\mathbf{X}, \mu, \tau) &= p(\mathbf{X} \mid \mu, \tau) p(\mu \mid \tau) p(\tau) = \\
&= \prod_{n=1}^N \mathcal{N}(x_n \mid \mu_N, \tau_N^{-1}) \mathcal{N}(\mu_N, (\lambda_N \tau_N)^{-1}) \text{Gamma}(a_N, b_N)
\end{aligned}$$

Is it correct, or should it be $\mathcal{N}(X \mid \mu_N, \tau_N^{-1}) \cdot \mathcal{N}(\mu_0, (\lambda_0 \tau_N)^{-1}) \cdot \text{Gamma}(a_0, b_0)$. How the hell I get the posterior!?

g

Chapter 2

Distributed Kalman filter

This chapter firstly describes basic (non-distributed) version of Kalman filter. This is later extended to a distributed variant.

2.1 Introduction

Kalman filter (KF) is widely used adaptive filtering method firstly described in *A New Approach to Linear Filtering and Prediction Problems* by R.E. Kalman in 1960. The filter is a recursive solution to the discrete-data linear filtering problem [16]. The filter is defined by a set of relatively simple equations and allows for efficient online computation without having to know the whole history of data and it is the optimal linear filter (under several conditions described below). It is one of the solution to one of the most fundamental problems in control theory, the linear quadratic Gaussian control problem. Over the years, numerous variants and extensions have been derived from the original Kalman filter and this thesis is concerned with one such extension – distributed Kalman filtering. See 2.6 for a brief overview of used methods.

Kalman filter has been thoroughly studied over the years and also adopted by multiple industries, having a vast amount of applications wherever an estimation of the state of process is needed and noise in prediction or observation is present. Because of its nature, it is commonly used in navigation and control of robots and vehicles[17], including spacecrafts[18]. It can be applied in time series analysis and hence is used in fields such as econometrics[19] and economics[20] or signal processing[21]. Various applications can be found in a medicine, such as neuroscience[22].

With respect to the current development of the so called Internet of Things and development of cheap measurements sensors and microcomputers, Distributed Kalman Filtering (KDF) is becoming even more relevant. As one of the solution to *sensor fusion* problem, it has been under an intensive research as can be seen in [23]. A network consisting of multiple sensors has many advantages over a single sensor scenario, such as robustness to noise, better field of view or, if designed correctly, not having a single point of failure [24]. Applications ranges from wireless networks[25], precision agriculture[Adamchuk2011], military and civil surveillance, medical applications, nuclear hazard assessment and others[26].

If needed, I could add information about the application - where is that predict-update part in each of them.

2.2 Description

The basic Kalman filter is used to estimate the current state of a linear dynamic system in case of noisy measurements. The state can be described by various variables, such as position and velocity. To use a Kalman Filter, one needs to specify system's:

1. control inputs, if any, such as sending a signal to steer a wheel
2. dynamic model, such as physical laws of motion
3. measurements, such as readings from sensors

The filter is useful in situations when dynamic model nor measurements cannot be entirely trusted (otherwise it one would not need to use it at all). It combines the information and while allowing for some noise – uncertainty – during prediction and measurements. The both pieces of information are weighted by so called Kalman's gain and depending on its value favors prediction over measurements and vice versa. The whole process is performed recursively, remembering the last estimate and covariance and hence does not need an entire history, which makes it attractive where memory is expensive. As a result a less noisy (hence a filter) and if used appropriately, better estimates are generated in each step, lying in between measurements and prediction.

Observable parameters coming from measurements does not have to be necessarily complete in a sense of describing the system in its entirety. It is a common case in which one observes only a subset of state variables and the filter can still estimate the entire state.

The distributed extension is easy to describe – instead of having only a single sensor, the system is measured by multiple of them and then somehow combining these measurements together. Despite its simple high level description, multisensor network of sensors must deal with various non-trivial problems and the whole computation is complicated based on the selected parameters of the network's behavior and topology. It also means getting into a distinctive research area of distributed data fusion[27] or consensus theory, which are out of the scope of this thesis.

obrázek:

<https://medium.com/@mithi/object-tracking-and-fusing-sensor-measurements-using-the-extended-kalman-filter-algorithm-part-1-f2158ef1e4f0>

<https://www.google.cz/search?q=kalman+filter&safe=off&source=lnms&tbm=isch&sa=X&ved=0ahUKEv>

2.3 Mathematical definition

Definition 3. Let $\mathbf{x}_k, \mathbf{x}_{k-1}, \mathbf{u}_k \in \mathbb{R}^n$ be the true state vectors in times $k, k-1$ and control vector, respectively. $\mathbf{z}_k \in \mathbb{R}^m$ is the measurement vector, $\mathbf{F}_k, \mathbf{B}_k \in \mathbb{R}^{n \times m}$ the state transition matrix and control matrix, respectively and $\mathbf{H}_k \in \mathbb{R}^{m \times n}$ is the observation matrix. $\mathbf{w}_k \in \mathbb{R}^n$ and $\mathbf{v}_k \in \mathbb{R}^m$ are process noise and observation noise, respectively. The filter assumes that the state evolves from (discrete) time $k-1$ to a new state in time k according to:

$$\mathbf{x}_k = \mathbf{F}_k \mathbf{x}_{k-1} + \mathbf{B}_k \mathbf{u}_k + \mathbf{w}_k \quad (2.1)$$

and then an (imperfect) observation is made

$$\mathbf{z}_k = \mathbf{H}_k \mathbf{x}_k + \mathbf{v}_k \quad (2.2)$$

In this work, the $\mathbf{B}_k \mathbf{u}_k$ part is omitted as is common in many scenarios, where control input does not make sense (such as temperature sensors). The process and observation noise $\mathbf{w}_k, \mathbf{v}_k$

2.5 Noise analysis

What is their purpose and what problems we have with them, why it's necessary to nicely estimate them otherwise everything breaks down (described in the main VBAKF article on the beginning) + optimality issue.

2.6 Methods

This is what I want: [23]

Survey of currently used methods to deal with unknown noise

2.7 Distributed Kalman Filter

Chapter 3

Distributed Variational Bayes Adaptive Kalman Filtration

This chapter describes main contributions of this work. It is a derivation of an improvement of a Variation Bayes Adaptive Kalman Filtration (VBAKF) described [30] by partially sharing parameters in diffusion networks from [26]. Up to the author knowledge, there has not been a similar attempt to combine these two, although [26] mention these as a future possibility.

We also propose a simplified version of the algorithm from [30] using exponential family form, as we believe that the resulting update equations are more elegant, concise and understandable.

3.1 Problem formulation

As has been described in

tady by měl být odkaz do teorie kde bude to co oni ají v Main Results.A + problémy s estimation šumu a proč se to musí řešit elegantně.

, our aim is ti describe a model in which we will infer not only \mathbf{x}_k , but also PNCM $\mathbf{P}_{k|k-1}$ and MCNM \mathbf{R}_k as opposed to traditional Kalman filter where these two are considered static

Popsat co ty zkratky znamenají

Based on explanation in

Odkaz, jak se ty rovnice přeloží přes HMM do pravděpodobnosti:

https://www.wikiwand.com/en/Kalman_filter#/Relationship_to_recursive_Bayesian_estimation

, the predicted PDF for \mathbf{x}_k and likelihood PDF

what is PDF

follows Gaussian distribution:

$$p(\mathbf{x}_k | \mathbf{z}_{1:k-1}, \mathbf{P}_{k|k-1}) = N(\mathbf{x}_k; \hat{\mathbf{x}}_{k|k-1}, \mathbf{P}_{k|k-1}) \quad (3.1)$$

$$p(\mathbf{z}_k | \mathbf{x}_k, \mathbf{R}_k) = N(\mathbf{z}_k; \mathbf{H}_k \mathbf{x}_k, \mathbf{R}_k) \quad (3.2)$$

where

- $N(\cdot; \mu, \Sigma)$ is Gaussian PDF with corresponding mean μ and covariance matrix Σ
- $\hat{\mathbf{x}}_{k|k-1}$ is the predicted state vector based on the model in

asi zase 2

- $\mathbf{P}_{k|k-1}$ is the predicted PNCM based on the model in

asi zase 2

We emphasize that $\hat{\mathbf{x}}_{k|k-1}$ and $\mathbf{P}_{k|k-1}$ are both inaccurate because precise \mathbf{Q}_k from the underlying update equation in

to samé co v předchozím todo

is unknown and we would not have to do the whole machinery if it was known.

3.1.1 Prior distributions for $\mathbf{P}_{k|k-1}$ and \mathbf{R}_k

We are going to leverage the benefits of good choice of priors to make the whole solution tractable outlined in

odkaz do jedničky?

. To ensure that the posterior belongs to the same distribution family as the prior distribution, we choose Inverse-Wishart distribution which is in a Bayesian statistic a common choice as a conjugate prior for the covariance matrix of a multivariate normal distribution with known mean[31]. That is the case for $\mathbf{P}_{k|k-1}$ and \mathbf{R}_k , which are both the covariances matrices of Gaussian PDFs as is in equations 3.1 and 3.2.

Definition 4. Inverse-Wishart distribution is a probability distribution on real-value positive-definite matrices. We say that \mathbf{X} follows the Inverse-Wishart distribution with ν degrees of freedom and scale matrix Ψ , i.e. $\mathbf{X} \sim \text{IW}(\Psi, \nu)$, if \mathbf{X}^{-1} has a Wishart distribution $\mathbf{X} \sim \text{W}(\Psi^{-1}, \nu)$. The density is

$$\frac{|\Psi|^{\frac{\nu}{2}}}{2^{\frac{\nu p}{2}} \Gamma_p(\frac{\nu}{2})} |\mathbf{X}|^{-\frac{\nu+p+1}{2}} e^{-\frac{1}{2}\text{tr}(\Psi \mathbf{X}^{-1})} \quad (3.3)$$

where $\Psi > \mathbf{0} \in \mathbb{R}^{p \times p}$ is a positive-definite scale matrix, $\nu > p - 1, \nu \in \mathbb{R}$ are degrees of freedom, $\Gamma_p(\cdot)$ is the multivariate gamma function and tr is the trace function. The exponential family form we are going to take advantage of in the following text is

$$\exp\left\{\begin{bmatrix} \mathbf{X}^{-1} \\ \ln |\mathbf{X}| \end{bmatrix} \begin{bmatrix} -\frac{1}{2}\Psi & -\frac{\nu+p+1}{2} \end{bmatrix} - \left(\frac{\nu}{2}(p \ln 2 - \ln |\Psi|) + \ln \Gamma_p\left(\frac{\nu}{2}\right)\right)\right\} \quad (3.4)$$

and the corresponding parts of the form from

tady odkaz na to kde definuju to exponential family form

are

nejsem si jistý jestli je to takhle správně - nemá být η a Ξ prohozené?

$$\begin{aligned} \eta(\mathbf{X}) &= \begin{bmatrix} \mathbf{X}^{-1} \\ \ln |\mathbf{X}| \end{bmatrix} \\ \Xi(\Psi, \nu, p) &= \begin{bmatrix} -\frac{1}{2}\Psi & -\frac{\nu+p+1}{2} \end{bmatrix} \\ A(\eta) &= \frac{\nu}{2}(p \ln 2 - \ln |\Psi|) - \ln \Gamma_p\left(\frac{\nu}{2}\right) \end{aligned}$$

The first element Ξ_1 of a *pseudovector* Ξ is a matrix, why the second Ξ_2 is a scalar. The advantages of this notation will be clearer when an addition to natural paramaters is carried over. The expected value of the IW distributed variable is

$$\mathbf{E}[\mathbf{X}] = \frac{\Psi}{\nu - p - 1} \quad (3.5)$$

$$= \frac{2}{2\Xi_2 + p + 1} \Xi_1 \quad (3.6)$$

Following from the above, the probabilities for $\mathbf{P}_{k|k-1}$ and \mathbf{R}_k are hence

$$p(\mathbf{P}_{k|k-1} | \mathbf{z}_{1:k-1}) = \text{IW}(\mathbf{P}_{k|k-1}; \hat{t}_{k|k-1}, \hat{\mathbf{T}}_{k|k-1}) \quad (3.7)$$

$$p(\mathbf{R}_k | \mathbf{z}_{1:k-1}) = \text{IW}(\mathbf{R}_k; \hat{u}_{k|k-1}, \hat{\mathbf{U}}_{k|k-1}) \quad (3.8)$$

Next, prior parameters for $\hat{t}_{k|k-1}, \hat{\mathbf{T}}_{k|k-1}, \hat{u}_{k|k-1}, \hat{\mathbf{U}}_{k|k-1}$ to represent the initial knowledge as precisely as possible. These parameters will be consumed later into prior's hyperparameters in Ξ .

Hyperparameters for PNCM

The model from [30] proposes the following initial parameter setting for the prior information for $\mathbf{P}_{k|k-1}$

$$\frac{\hat{\mathbf{T}}_{k|k-1}}{\hat{t}_{k|k-1} - n - 1} = \tilde{\mathbf{P}}_{k|k-1} = \mathbf{F}_{k-1} \mathbf{P}_{k-1|k-1} \mathbf{F}_{k-1}^T + \tilde{\mathbf{Q}}_{k-1} \quad (3.9)$$

$$\hat{\mathbf{T}}_{k|k-1} = \tau \tilde{\mathbf{P}}_{k|k-1} \quad (3.10)$$

$$\hat{t}_{k|k-1} = n + \tau + 1 \quad (3.11)$$

where $\tilde{\mathbf{P}}_{k|k-1}, \tilde{\mathbf{Q}}_{k-1}$ are the *nominal* PECM and PNCM, respectively. The latter, together with $\tau \geq 0$, are model's parameters and their choice is discussed later. Substituting $\mathbf{P}_{k|k-1}, \hat{t}_{k|k-1}, \hat{\mathbf{T}}_{k|k-1}$ into 3.4 and combining with 3.10 and 3.11, we get

$$\eta(\mathbf{P}_{k|k-1}) = \begin{bmatrix} \mathbf{P}_{k|k-1}^{-1} \\ \ln |\mathbf{P}_{k|k-1}| \end{bmatrix} \quad (3.12)$$

$$\Xi(\hat{\mathbf{T}}_{k|k-1}, \hat{t}_{k|k-1}, n) = \begin{bmatrix} -\frac{1}{2} \hat{\mathbf{T}}_{k|k-1} & -\frac{\hat{t}_{k|k-1} + n + 1}{2} \end{bmatrix} = \quad (3.13)$$

$$= \begin{bmatrix} -\frac{1}{2} \tau \tilde{\mathbf{P}}_{k|k-1} & -\frac{1}{2} \tau - 1 - n \end{bmatrix} \quad (3.14)$$

where n is the dimension of $\hat{\mathbf{T}}_{k|k-1}$.

Hyperparameters for MCNM

According to 1.1, the prior for $p(\mathbf{R}_k | \mathbf{z}_{1:k-1})$ is

$$p(\mathbf{R}_k | \mathbf{z}_{1:k-1}) = \int p(\mathbf{R}_k | \mathbf{R}_{k-1}) p(\mathbf{R}_{k-1} | \mathbf{z}_{1:k-1}) d\mathbf{R}_{k-1} \quad (3.15)$$

To make sure that $p(\mathbf{R}_k | \mathbf{z}_{1:k-1})$ is also Inverse-Wishart distributed, the $p(\mathbf{R}_k | \mathbf{R}_{k-1})$ must have a suitable form so the Bayesian update results in the exponential family. Allowing slowly varying MNCM, as is the case in many real world applications, a simple heuristic from [32] is

chosen to deal with this uncertainty. That is spreading previous posteriors by a factor of ρ , which is then a model parameter

$$\hat{u}_{k|k-1} = \rho(\hat{u}_{k-1|k-1} - m - 1) + m + 1 \quad (3.16)$$

$$\hat{\mathbf{U}}_{k|k-1} = \rho \hat{\mathbf{U}}_{k-1|k-1} \quad (3.17)$$

where m is the size of matrix $\hat{\mathbf{U}}_{\cdot|\cdot}$. The initial MNCM \mathbf{R}_0 is also considered to be Inverse-Wishart distributed

$$p(\mathbf{R}_0 = \text{IW}(\mathbf{R}_0; \hat{u}_{0|0}, \hat{\mathbf{U}}_{0|0})) \quad (3.18)$$

Similarly to the PNCM case in 3.9, the initial prior information for the mean value of \mathbf{R}_0 is set as the initial nominal MNCM $\tilde{\mathbf{R}}_0$

$$\frac{\hat{\mathbf{U}}_{0|0}}{\hat{u}_{0|0} - m - 1} = \tilde{\mathbf{R}}_0 \quad (3.19)$$

where $\tilde{\mathbf{R}}_0$ is then a model parameter.

Substituting \mathbf{R}_{k-1} , \hat{u}_k , $\hat{\mathbf{U}}_k$ into 3.4 and combining with 3.17 and 3.16, we get

$$\eta(\mathbf{R}_{k-1}) = \begin{bmatrix} \mathbf{R}_{k-1}^{-1} \\ \ln |\mathbf{R}_{k-1}| \end{bmatrix} \quad (3.20)$$

$$\Omega(\hat{\mathbf{U}}_k, \hat{u}_k, m) = \begin{bmatrix} -\frac{1}{2}\hat{\mathbf{U}}_k & -\frac{\hat{u}_k + m + 1}{2} \end{bmatrix} = \quad (3.21)$$

$$= \begin{bmatrix} -\frac{1}{2}\rho\hat{\mathbf{U}}_k & -\frac{1}{2}\rho(\hat{u}_k - m - 1) - m - 1 \end{bmatrix} \quad (3.22)$$

where n is the dimension of $\hat{\mathbf{U}}_k$.

Choice of the nominal PNCM and MNCM

In the detailed analysis [30] can be found that to guarantee model's convergence, the initial nominal PNCM $\tilde{\mathbf{Q}}_k$ needs to be near the initial true PNCM \mathbf{Q}_k . In the paper and as well in this work, we set $\mathbf{Q}_k = \text{diag}[\alpha_{1,k}, \dots, \alpha_{n,k}]$, $\alpha_{i,k} > 0$. In reality, this is done based on the state-of-the-art engineering knowledge in the domain. The same is applied for the initial nominal MNCM $\tilde{\mathbf{R}}_0$ and the choice is $\mathbf{R}_k = \text{diag}[\beta_1, \dots, \beta_m]$, $\beta_i > 0$ and the reader is kindly redirected to the original study for more information.

3.2 Variational Bayes

In order to infer all \mathbf{x}_k , $\mathbf{P}_{k|k-1}$ and \mathbf{R}_k , a posterior PDF $p(\mathbf{x}_k, \mathbf{P}_{k|k-1}, \mathbf{R}_k)$ needs to be estimated. As thoroughly described in chapter 1.7, the Variation Bayes in our case expects the following approximation

$$p(\mathbf{x}_k, \mathbf{P}_{k|k-1}, \mathbf{R}_k | \mathbf{z}_{1:k}) \approx q(\mathbf{x}_k)q(\mathbf{P}_{k|k-1})q(\mathbf{R}_k) \quad (3.23)$$

where $q(\mathbf{x}_k)$, $q(\mathbf{P}_{k|k-1})$, $q(\mathbf{R}_k)$ are factors of our factorized distribution $q(\cdot)$ by which we approximate the $p(\cdot)$. These factors are obtained using Variational Bayes machinery including minimizing the Kullback-Leibler divergence and using CAVI.




Figure 3.1: Graphical model of probabilistic model 3.24

3.2.1 Probabilistic model

Combining 3.1, 3.2, 3.7 and 3.8, the join PDF can be factored as

je tady důvod proč oni raději použili $p(\mathbf{z}_{1:k})$?

$$p(\boldsymbol{\Xi}, \mathbf{z}_{1:k}) = p(\mathbf{z}_k | \mathbf{x}_k, \mathbf{R}_k) p(\mathbf{x}_k | \mathbf{z}_{1:k-1}, \mathbf{P}_{k|k-1}) p(\mathbf{P}_{k|k-1} | \mathbf{z}_{1:k-1}) p(\mathbf{R}_k | \mathbf{z}_{1:k-1}) p(\mathbf{z}_{1:k-1}) \quad (3.24)$$

and substituting appropriate distribution for these PDFs:

$$p(\boldsymbol{\Xi}, \mathbf{z}_{1:k}) = N(\mathbf{z}_k; \mathbf{H}_k \mathbf{x}_k, \mathbf{R}_k) N(\mathbf{x}_k; \hat{\mathbf{x}}_{k|k-1}, \mathbf{P}_{k|k-1}) \quad (3.25)$$

$$\times \text{IW}(\mathbf{P}_{k|k-1}; \hat{t}_{k|k-1}, \hat{\mathbf{T}}_{k|k-1}) \text{IW}(\mathbf{R}_k; \hat{u}_{k|k-1}, \hat{\mathbf{U}}_{k|k-1}) \quad (3.26)$$

$$\times p(\mathbf{z}_{1:k-1}). \quad (3.27)$$

graphical model and exponential variant

3.2.2 Deriving update equations

As opposed to [30], we are going to leverage advantages of exponential family forms to infer the update equations. Considering the exponential form of a Gaussian distribution

$$N(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \exp \left\{ \begin{bmatrix} \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu} \\ -\frac{1}{2} \boldsymbol{\Sigma}^{-1} \end{bmatrix} \begin{bmatrix} \mathbf{x} & \mathbf{x} \mathbf{x}^T \end{bmatrix} - \left(\frac{1}{2} \boldsymbol{\mu}^T \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu} + \frac{1}{2} \ln |\boldsymbol{\Sigma}| \right) \right\} \quad (3.28)$$

$$= \exp \left\{ \begin{bmatrix} \boldsymbol{\Sigma}^{-1} \\ \ln |\boldsymbol{\Sigma}^{-1}| \end{bmatrix} \begin{bmatrix} -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T & -\frac{1}{2} \end{bmatrix} + \ln(2\pi) \right\} \quad (3.29)$$

where $\mathbf{x}\mathbf{x}^T$ is an outer product, hence resulting into a matrix from $\mathbb{R}^{n \times n}$. Substituting appropriate factors for both \mathbf{x}_k as well as \mathbf{z}_k we get

$$N(\mathbf{x}_k; \hat{\mathbf{x}}_{k|k-1}, \mathbf{P}_{k|k-1}) = \exp \left\{ \left[\begin{array}{c} \mathbf{P}_{k|k-1}^{-1} \\ \ln |\mathbf{P}_{k|k-1}^{-1}| \end{array} \right] \left[\begin{array}{cc} -\frac{1}{2}(\mathbf{x}_k - \hat{\mathbf{x}}_{k|k-1})(\mathbf{x}_k - \hat{\mathbf{x}}_{k|k-1})^T & -\frac{1}{2} \end{array} \right] - k_1 \right\} \quad (3.30)$$

$$N(\mathbf{z}_k; \mathbf{x}_k, \mathbf{R}_k) = \exp \left\{ \left[\begin{array}{c} \mathbf{R}_k^{-1} \\ \ln |\mathbf{R}_k^{-1}| \end{array} \right] \left[\begin{array}{cc} -\frac{1}{2}(\mathbf{z}_k - \mathbf{x}_k)(\mathbf{z}_k - \mathbf{x}_k)^T & -\frac{1}{2} \end{array} \right] - k_2 \right\} \quad (3.31)$$

Důležité: Tady mi nesedí ty updates!

As can be seen, the natural parameters matches for 3.30 and 3.12 as well as for 3.31 and 3.20, which is of course intentional and result of clever choice of conjugated priors.

With respect to

odkaz na mr. Dedeciovo tvrzení s $T+=1$, ...

, the single step Bayesian update of Inverse-Wishart hyperparameters with Multivariate Normal Distribution can be written as

$$\Xi_i \leftarrow \underbrace{\left[\begin{array}{c} -\frac{1}{2}\Psi \\ -\frac{\nu + p + 1}{2} \end{array} \right]}_{\Xi_{i-1}} + \underbrace{\left[\begin{array}{c} -\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T \\ -\frac{1}{2} \end{array} \right]}_{T(x, \mu)} \quad (3.32)$$

Utilizing the fact that

TADY U TOHOHLE NEMŮŽU BÝT VĚTŠÍ ALIBISTA. Nerozumím tomu, jak to dostali a ani proč je to vlastně potřeba. Navíc nevím, jak tam dostat tu střední hodnotu. Je mi jasné, že to nějak vyleze z VB kdy je tam $E[\log(p)]$, ale nevidím zkrz to.

$$E^i \left[-\frac{1}{2}(\mathbf{x}_k - \hat{\mathbf{x}}_{k|k-1})(\mathbf{x}_k - \hat{\mathbf{x}}_{k|k-1})^T \right] = \mathbf{P}_{k|k}^i + (\hat{\mathbf{x}}_{k|k}^i - \hat{\mathbf{x}}_{k|k-1})(\hat{\mathbf{x}}_{k|k}^i - \hat{\mathbf{x}}_{k|k-1})^T \quad (3.33)$$

$$E^i \left[-\frac{1}{2}(\mathbf{z}_k - \mathbf{H}_k \mathbf{x}_k)(\mathbf{z}_k - \mathbf{H}_k \mathbf{x}_k)^T \right] = \mathbf{H}_k \mathbf{P}_{k|k}^i \mathbf{H}_k^T + (\mathbf{z}_k - \mathbf{H}_k \hat{\mathbf{x}}_{k|k}^i)(\mathbf{z}_k - \mathbf{H}_k \hat{\mathbf{x}}_{k|k}^i)^T \quad (3.34)$$

the single step i of the variational updates are

$$\Xi_{i+1}^n \leftarrow \Xi^{n-1} + \left[\begin{array}{c} -\frac{1}{2}(\mathbf{P}_{k|k}^i + (\hat{\mathbf{x}}_{k|k}^i - \hat{\mathbf{x}}_{k|k-1})(\hat{\mathbf{x}}_{k|k}^i - \hat{\mathbf{x}}_{k|k-1})^T) \\ -\frac{1}{2} \end{array} \right] \quad (3.35)$$

$$\Omega_{i+1}^n \leftarrow \Omega^{n-1} + \left[\begin{array}{c} \mathbf{H}_k \mathbf{P}_{k|k}^i \mathbf{H}_k^T + (\mathbf{z}_k - \mathbf{H}_k \hat{\mathbf{x}}_{k|k}^i)(\mathbf{z}_k - \mathbf{H}_k \hat{\mathbf{x}}_{k|k}^i)^T \\ -\frac{1}{2} \end{array} \right] \quad (3.36)$$

where Ξ^{n-1}, Ω^{n-1} are hyperparameters obtained from the results of Kalman filter of the previous observation, derived using formulas 3.14 and 3.22.

Kalman correction

The last bit of the single CAVI iteration is the Kalman correction, which is the standard

$$\mathbf{K}_k^{(i+1)} = \hat{\mathbf{P}}_{k|k-1}^{(i+1)} \mathbf{H}_k^T (\mathbf{H}_k \hat{\mathbf{P}}_{k|k-1}^{(i+1)} \mathbf{H}_k^T + \hat{\mathbf{R}}_k^{(i+1)})^{-1} \quad (3.37)$$

$$\hat{\mathbf{x}}_{k|k}^{(i+1)} = \hat{\mathbf{x}}_{k|k-1} + \mathbf{K}_k^{(i+1)} (\mathbf{z}_k - \mathbf{H}_k \hat{\mathbf{x}}_{k|k-1}) \quad (3.38)$$

$$\mathbf{P}_{k|k}^{(i+1)} = \hat{\mathbf{P}}_{k|k-1}^{(i+1)} - \mathbf{K}_k^{(i+1)} \mathbf{H}_k \hat{\mathbf{P}}_{k|k-1}^{(i+1)} \quad (3.39)$$

3.2.3 Parameter diffusion

Thanks to using the exponentially family forms, parameter diffusion can be easily incorporated into the final algorithm – the nodes only need to exchange the hyperparameters Ξ^n and Ω^n . While the previous chapters were mostly based on algorithm described in [30], the parameter diffusion part is mostly based on [26]. We just present here the result algorithm, but the reader is advised to reach the paper for more details.

Consider a network of connected nodes, each of them doing their own observation. A single node i obtain a measurement in time t and proceed with Kalman filtration as above. After obtaining final $\mathbf{P}_{k|k}^{i,t}, \hat{\mathbf{x}}_{k|k}^{i,t}$ and also result hyperparameters $\Xi^{i,t}, \Omega^{i,t}$. The node then *contacts* its neighbor nodes $1, \dots, d$ and attempts to gather all available hyperparameters $\Xi^{j,t}, \Omega^{j,t}, j \in \{1, \dots, d\}$ based on the current conditions (e.g. only a single hop distance, or more). The node then performs a parameter fusion of these hyperparameters based on a chosen strategy represented by some function $\Xi_*^{i,t} = g(\Xi^{i,t}, \Xi^{1,t}, \dots, \Xi^{d,t})$. An example of such function may be just a simple averaging over all hyperparameters $g(\Xi^{i,t}, \Xi^{1,t}, \dots, \Xi^{d,t}) = \frac{\sum_{s=\{i,1,\dots,d\}} \Xi^{s,t}}{d+1}$. The situation is for the hyperparameter $\Omega^{i,t}$ is analogical. The node then sends this updated hyperparameters $\Xi_*^{i,t}, \Omega_*^{i,t}$ back to the cooperating nodes, which incorporates this hyperparameter into their knowledge.

3.3 Algorithm summary

Here we summarize the previous pages into a single algorithm. The single step of the proposed algorithm for a single node is then

Algorithm 5. 1. *Inputs:* $\hat{\mathbf{x}}_{k-1|k-1}, \mathbf{P}_{k-1|k-1}, \hat{u}_{k-1|k-1}, \hat{\mathbf{U}}_{k-1|k-1}, \mathbf{F}_{k-1}, \mathbf{H}_k, \mathbf{z}_k, \tilde{\mathbf{Q}}_{k-1}, m, n, \tau, \rho, N$
2. *Time Update of the Kalman Filter*

jen odkaz jako jinde

$$\begin{aligned} \hat{\mathbf{x}}_{k|k-1} &= \mathbf{F}_{k-1} \hat{\mathbf{x}}_{k-1|k-1} \\ \tilde{\mathbf{P}}_{k|k-1} &= \mathbf{F}_{k-1} \mathbf{P}_{k-1|k-1} \mathbf{F}_{k-1}^T + \tilde{\mathbf{Q}}_{k-1} \end{aligned}$$

3. Initialize Ξ_{init} by 3.14, Ω_{init} by 3.22, $\hat{\mathbf{x}}_{k|k}^{(0)} = \hat{\mathbf{x}}_{k|k-1}, \mathbf{P}_{k|k}^{(0)} = \tilde{\mathbf{P}}_k$

For i **in** N **do:**

4. Get new Ω^i based on the 3.36 and compute the $E^i[\mathbf{R}_k]$ 3.6

5. Get new Ξ^i based on the 3.35 and compute $E^i[\mathbf{P}_{k|k-1}]$ using 3.6

6. Perform Kalman correction 3.37-3.39 using expected values to obtain corrected $\hat{\mathbf{x}}_{k|k}^{i+1}, \mathbf{P}_{k|k}^{i+1}$

end for

7. *Perform parameter diffusion, i.e. if hyperparameters from nodes $j \in \{1, \dots, d\}$ are available, get*

$$\Xi_*^{i,t} = g(\Xi^{i,t}, \Xi^{1,t}, \dots, \Xi^{d,t}) \quad (3.40)$$

$$\Omega_*^{i,t} = g(\Omega^{i,t}, \Omega^{1,t}, \dots, \Omega^{d,t}) \quad (3.41)$$

8. **Returns:** $\hat{\mathbf{x}}_{k|k}^{i+1}, \mathbf{P}_{k|k-1}^{i+1}, \Xi_*^{i,t}, \Omega_*^{i,t}$

Chapter 4

Results

The following chapter offers an in-depth analysis of a performance of the proposed algorithm. To guarantee a reproducible research, implementation details are mentioned to some degree and the code with the comparison is publicly available on

add link to the github page

. This thesis focused mostly on state prediction accuracy and less on the noise analysis, although it is mentioned when it helps to understand the mechanism behind the performance.

Please note that some figures should be compared against each other as a trend, rather than individual lines. This is true especially in the case of the multi-node variants, when the author's goal is to demonstrate that all the nodes are having some specific property which would be hardly to demonstrate by plotting only a single node.

4.1 Implementation

The implementation used in this work has been written in feature rich Python language ecosystem. Next to the CPython interpreter 3.6, following libraries (all licensed under MIT license) has been reused:

- scipy 1.0.0 – generator of random numbers from various distributions
- numpy 1.14.2 – arrays and mathematical operations, linear algebra
- networkx 2.1 – to create a network topology of measurement nodes
- jupyter 5.2.3 – “notebook” ecosystem for a rapid analysis
- filterpy 1.2.1 - reference implementation of a classic Kalman Filter
- matplotlib 2.2.2 – to create charts (including this thesis)

All the code is available on GitHub page

Add link to the github page

including the instruction how to replicate the analysis and testing environment. Most of the functionality has been written from scratch by the author while having a reference implementation by the supervisor (plus the model of the testing trajectory). The code is structured into several parts, notably objects `MeasurementNode` and `IWPrior` plus many of utility functions in the appropriate modules such as `exponential_family.py` or `network.py`, which are self-explanatory.

Comparison

The notation for the components of the state vector is kept as x_i in the charts. From the nature of the problem under study, it is clear that first two components which are observable behave in the similar way and the same is true for the latter two (which are not observable). To keep the analysis clean and concise, only a representative component is plotted if the behavior is significantly different.

The main metric used for comparison is root mean square error. Where appropriate, standard deviation of errors is listed. The result numbers discussed further are results of many iterations (more than 20) with different initial noise assigned to different nodes. It was although made sure that when comparing two different variants, same noise has been attributed to the appropriate elements. Finally, in situations when we are going to *visualize* matrices behavior, in our case matrices close to a matrix $c\mathbf{I}, c > 0$, a Frobenius matrix norm is used.

4.2 Simulations

4.2.1 Testing data

Here we describe the dynamical system model in the standard notation suitable for the Kalman Filter. The state transition matrix was of the form

$$\mathbf{F} = \begin{bmatrix} 1 & 0 & t & 0 \\ 0 & 1 & 0 & t \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}, \mathbf{Q} = 2 \cdot \begin{bmatrix} t^3/3 & 0 & t^2/2 & 0 \\ 0 & t^3/3 & 0 & t^2/2 \\ t^2/2 & 0 & t & 0 \\ 0 & t^2/2 & 0 & t \end{bmatrix}, t = 0.1$$

$$\mathbf{H} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}, \mathbf{R} = 0.25 \cdot \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

A random trajectory was then simulated for all four components using this model with multivariate Gaussian random noise added to the each “true” datapoint (with covariance \mathbf{Q}) as well as to the observations (with covariance \mathbf{R}). It can be thought about as a point moving on a 2D plane. A typical result of the first two components of the test trajectory can be seen on 4.1 while the other two on 4.2. We always generated 250 time steps.

Noise

Different type of noises were tested to thoroughly examine the behavior of different filter under different circumstances. Examples for a static noise and variable noise can be found on figures 4.3 and 4.4. Even though only a single line is plotted, in the simulation both observable components distorted by the type of noise under study. Henceforth, for the sake of simplicity, MNCM was the factor we considered that can be slowly varying.

It worth mentioning that because of the task’s triviality, the CAVI algorithm usually converges in few iterations – less than 5. Therefore, number of iterations were set to conservative 10. There is an exhaustive analysis of algorithm stability and parameters effect in [30], where the reader is kindly referred to if she wishes to know more.




Figure 4.1: First and second component in the first row are observable, while the third and fourth is not.




Figure 4.2: True third and fourth component of the state vector.

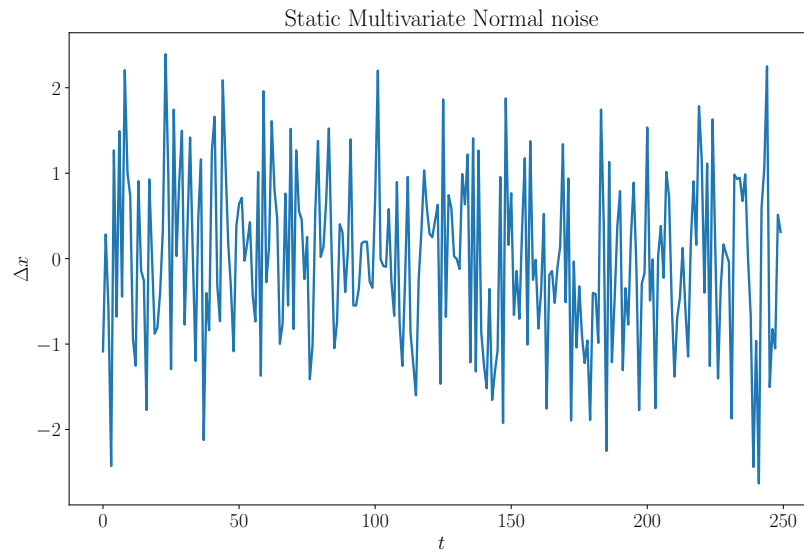


Figure 4.3: Multivariate Gaussian Random noise without variable variance.




Figure 4.4: A variable variance in the noise is simulated in this example.

	x_1	x_2	x_3	x_4
VBAKF	0.65	0.578	0.96	1.045
CKF	0.601	0.555	0.937	0.935

Table 4.1: Comparison of RMSE of a single node under static noise when the noise covariance is known.

4.2.2 Single node VBAKF

In this section, we compare the VBAKF algorithm performance against the classic Kalman Filter (CKF). It is important to distinguish between the situation when the true noise covariance is known and when it is not. Another orthogonal property of the system is the variance of the noise covariance over time. Since the motivation of developing VBAKF is specially for the situations when the covariance is unknown or there is a risk of choosing a wrong one, it is clear where the VBAKF is expected to perform better.

When the true covariance is known and static – an unlikely situation in the real world systems – we do not expect any improvements in predictions over VBAKF. Actually, its performance may be degraded trying to adjust to the noise, making the overall accuracy score worse than CKF. In the case when the covariance is known and changing, then it depends on how is one able to tune the algorithm parameters (such as ρ or τ), how big the changes are and how precise is the known value.

dát do conclusion, že VBAKF si ani tehdy nevedl moc špatně

Henceforth, the main advantage is in the case when one do not know the true values or we do not know the initial estimate (or both), as is often the case in reality. It is the scenario where we expect VBAKF to surpass CKF, since if we are at least somewhat close to the initial estimate, it should *correct* to more precise values over time. Additionally, if the covariance is slowly changing, than VBAKF is able to adjust it as well.

Static noise

As is clear from the table 4.1 , the claim mentioned in the previous section is validated. The CKF variant is slightly outperforming VBAKF, although both variants are performing quite well. Figures 4.5 and 4.6 allows for some introspection to the observed behavior. VBAKF algorithm seems to pay the price for being able more flexible than CKF and it is following noise peaks too much. This could be to some degree controlled by the τ hyperparameter, but not much without degrading performance in some other means, such as being to rigid, or coalesce with CKF.

Let also demonstrate the situation when the true covariance is not known and set to some reasonable estimate close to the true covariance. The remarkable property of being able to adjust to the true covariance has been confirmed, as can be seen in 4.7 , the performance is significantly better, specially in the predicted-only components of the state vector. This is even more interesting when considering that the standard deviation of prediction error summarized in the 4.8 is significantly lower for VBAKF than in CKF. This is the remarkable property of the proposed algorithm and figures for the first component on figure 4.9 and third component on 4.10 only supports that. Both algorithms started with noise covariance estimate which was off, but VBAKF variant could adjust for it, placing less emphasis on the measurements and believing more to the prediction step. CKF algorithm just stays on the same level of reliance with correction.




Figure 4.5: Performance of both filters for x_1 together with measurements and the true trajectory.




Figure 4.6: Performance of both filters for x_3 together with the true trajectory. This is predicted variable and not directly observed, hence the measurements are missing.




Figure 4.7: RMSE comparison of both algorithms over components.

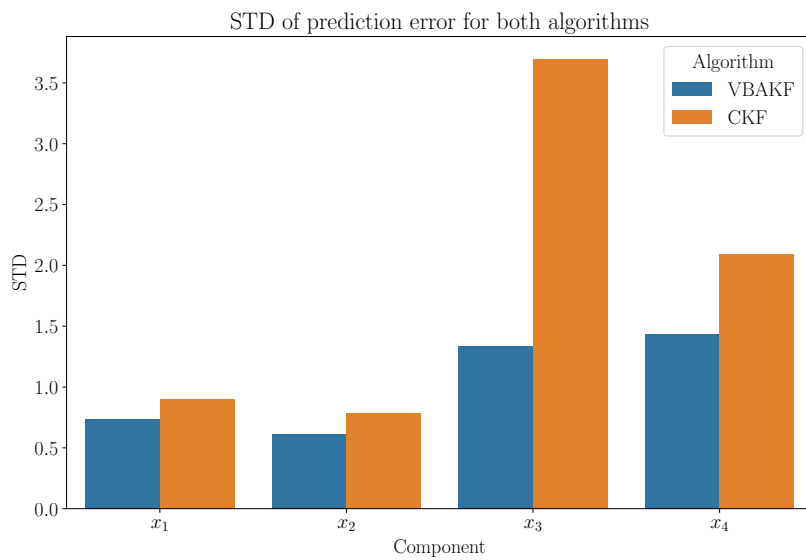


Figure 4.8: Standard deviation of the prediction error across components.




Figure 4.9: First and second component in the first row are observable, while the third and fourth is not.




Figure 4.10: Predicted values of x_3 component by both algorithms and with true trajectory.




Figure 4.11: First component estimates for both algorithms with observations, true values and also the noise amplitude. Notice how VBAKF is becoming more rigid during high noise peaks while still being flexible to catch up to the trajectory change.

Variable noise

In this scenario, it is again considered that the value for true covariance is unknown and contrary to the previously discussed static case, it is slowly changing as is visualized on the figure 4.4. As can be seen from trajectories on figures 4.11 and 4.12, VBAKF is again able to put less emphasis on the measurements during epochs of high noise, contrary to CKF being too sensitive in during these.

Finally, a very nice visualization of VBAKF capabilities can be seen on the graph 4.13. The adaptation of the noise covariance matrix to the true unknown covariance matrix is supporting the earlier findings of being able to account for noisy epochs.

Conclusion

nechat si do finálného conclusion?

The simulations shows that our single node VBAKF implementation is behaving as the theory predicts and in concordance with the reference source [30] where more emphasis on noise analysis were done. VBAKF does not perform better in case of a static noise when the covariance is known, merely just adding more computational time and possibly worse results trying to adjust to random fluctuations. It outperforms CKF in situations when the covariance of the noise is either unknown and, or, slowly changing. The performance of VBAKF was consistently better than CKF in these over many iterations and many variant of a noise. It is also worth mentioning that a very little tuning had to be done to get to these results. How much close the initial estimate must be to guarantee divergence is of course almost impossible to say in the real world cases, but in our testing case, even extreme initialization lead to quickly correcting the state (in less than 5 iterations). Such an extreme initialization were for example over three orders of magnitudes of the reality – a situation in which the CKF just does not provide any useful information.




Figure 4.12: Third component estimates for both algorithms with observations, true values and also the noise amplitude. As in the case of 4.11, the VBAKF can control during very noisy intervals.




Figure 4.13: Frobenius norms comparison of the VBAKF algorithm and covariance of moving window (size 20 steps to the future) of the noise. The diagonal values of \mathbf{R} are increasing during noise peaks, while going down during low noise epochs. The scale in this chart is not important, it is the trend of the norms. Those should be ideally as close as possible.




Figure 4.14: A network representing connection of different measurement nodes connected with the neighbors. The topology is a 4-regular graph.

4.2.3 Distributed VBAKF

Based on the results from the previous chapter 4.2.2, only a slowly changing noise is going to be considered to not clutter the analysis of VBAKF with diffusion versus VBAKF without diffusion. As has been mentioned on the beginning of this whole chapter, the figures when plotting information from multiple nodes should be read as a general trend exhibited by the system. It would be a mistake try to interpret them line by line.

Setting

Our network topology of nodes is for simplicity a 4-regular graph as is in figure . A slowly changing, different noise is attributed to the each node in the network (figure). When we compare diffused versus non-diffused variant, the noise attributed in the both runs was the same in corresponding nodes to make the results comparable and reproducible. The results are averaged over at least 100 iterations.

Estimation performance

As expected, the estimation performance of the diffused variant is better than non-diffused variant. This is not only a result in the averaged results on the figure 4.16 , but it was consistent across literally all runs. To recapitulate – there was not a *single* run where would the non-diffused variant outperformed the diffused one in matter of RMSE and standard deviation (figure 4.17) of error. There is a price in time execution, since the diffusion variant does an additional step – diffusing parameters – in each iteration, but that was negligible, taking only about 5% longer than non-diffused variant.

Noise covariance analysis

The idea behind diffusion optimization is that the *averaging* over several nodes will eliminate extremes and smooths the estimation. This is clearly happening, as is clear from 4.18 , where we




Figure 4.15: Values of first and second component with noise from different nodes.

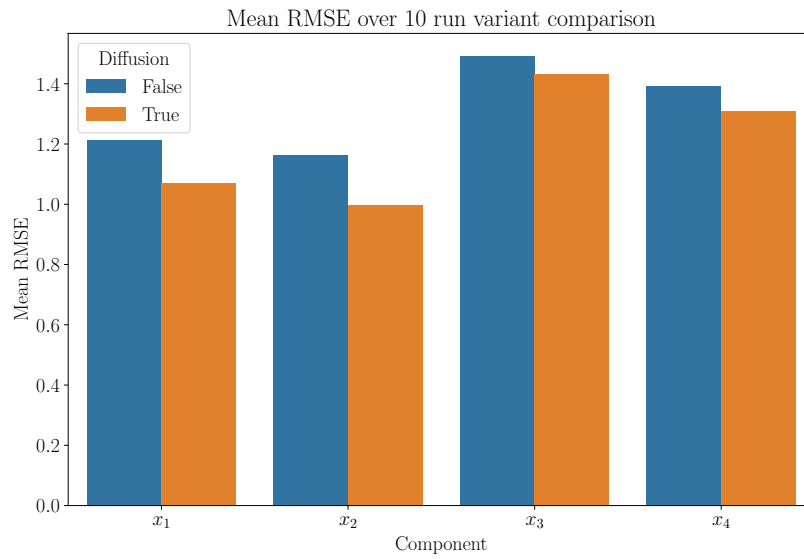


Figure 4.16: First and second component in the first row are observable, while the third and fourth is not.

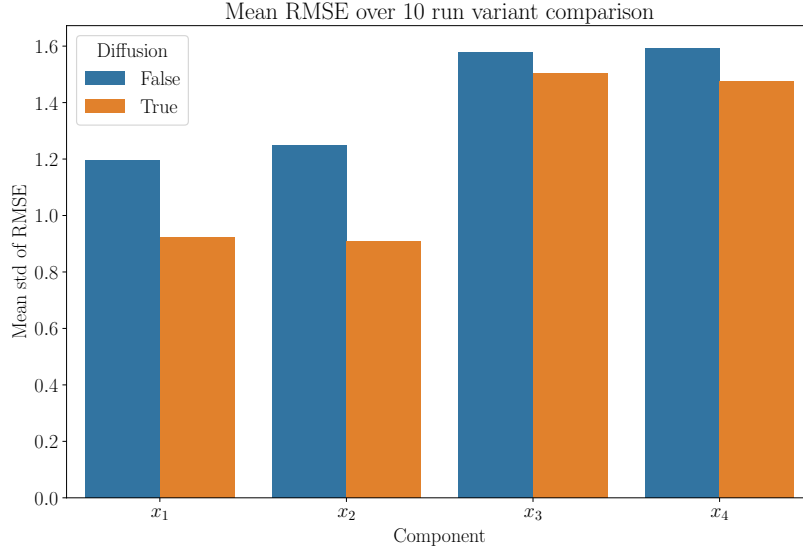


Figure 4.17: First and second component in the first row are observable, while the third and fourth is not.




Figure 4.18: The diffusion algorithm smooths the trajectory for this single randomly chosen single node from the network, ignoring significant peaks which would otherwise led to inferior estimates.

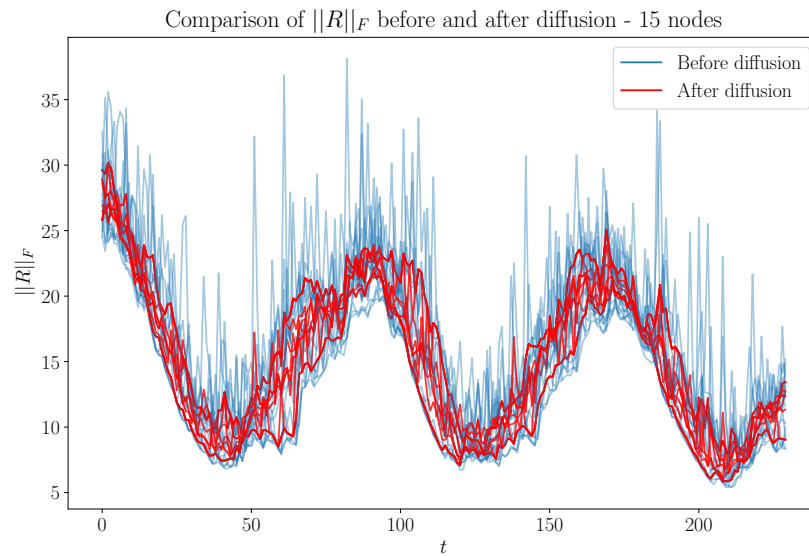


Figure 4.19: All nodes exhibit the behavior of smoothing out the noise covariance estimates, leading to the better results.

plot a Frobenius norm of the noise covariance matrix before and after diffusion step. Naturally, this phenomenon was exhibited on all nodes in the network, visualized on 4.19 .

Another view of this is presented on figures 4.20 and 4.21 . While there is almost no or very little pattern in the non-diffused variant, the diffusion variant got a general trend.

It is also interesting to note that in our simulations, the diffused variance exhibit much less divergence cases, which can be explained by the fact that any extremes are much likely to be suppressed by neighbor nodes participating in diffusion. An example can be seen on 4.22 .

hodit obrázky do dvou sloupců aby to nebylo tak rozlezlý jako teď. Určitě by tak mohli být udělány: příklady šumů, bar plot RMSE a STD.

optional: fig: RMSE ve dvou složkách v závislosti na počtu nodes pro obě varianty * tohle nepůjde - hází to různá čísla a bude to dle mého šumem - neřešit a jet zpětná vazba

Conclusion a nebo jen finální conclusion




Figure 4.20: Nodes in non-diffusion variant do not exhibit any common pattern. There is also one of the nodes diverging quite significantly, considering the logarithmic scale. The diffusion variant do not exhibit such divergence.

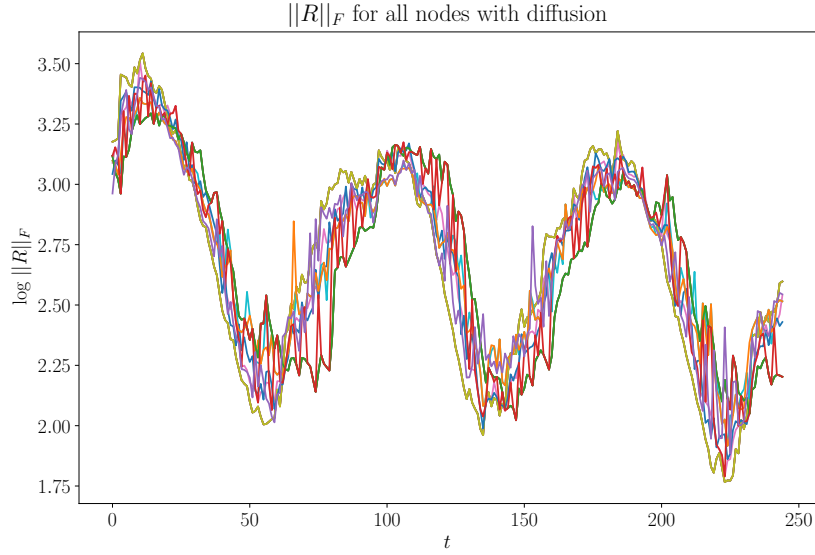


Figure 4.21: Nodes in the diffusion variant are able to pick up the general noise trend and “agree” on it to some degree, even though it may not be the best estimate for the underlying nodes.

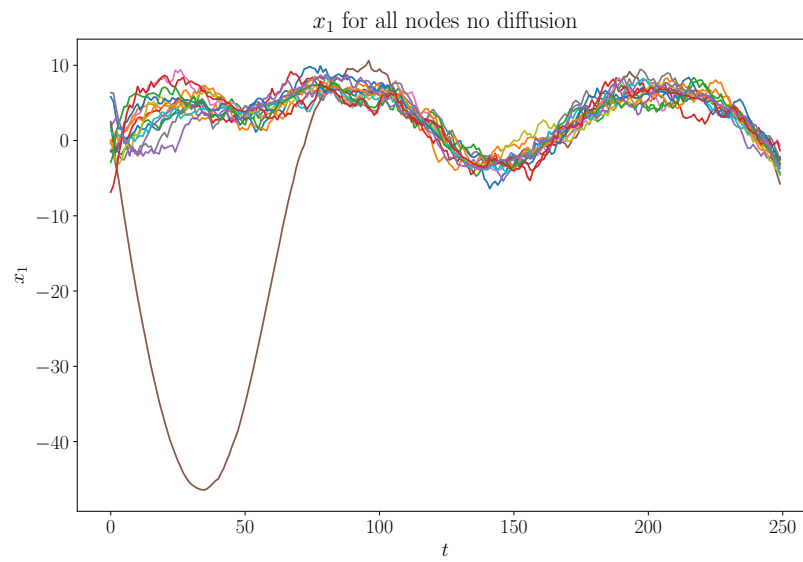


Figure 4.22: An example of diverging node present in the non-diffused variant.

Conclusion

Conclusion here

Bibliography

- [1] E. T. Jaynes. “Probability Theory: The Logic of Science.” In: *The Mathematical Intelligencer* 27.2 (2003), pp. 83–83. ISSN: 0343-6993. DOI: 10.1007/BF02985800. arXiv: arXiv:1011.1669v3.
- [2] Edward N ed Hájek Zalta. “Interpretations of Probability”. In: *The Stanford Encyclopedia of Philosophy*. 2012. ISBN: 1095-5054. URL: <https://plato.stanford.edu/>.
- [3] Maurice J Dupré and Frank J Tipler. “New Axioms for Rigorous Bayesian Probability”. In: *Bayesian Analysis* 1.1 (2004), pp. 1–8. URL: <http://dauns.math.tulane.edu/%7B%7Ddupre/OLDPUBLIC/20090804%20BA%20FNL%203.pdf>.
- [4] Charles Friedman. “The Frequency Interpretation in Probability”. In: *Advances in Applied Mathematics* 23.3 (1999), pp. 234–254. ISSN: 0196-8858. DOI: <https://doi.org/10.1006/aama.1999.0653>. URL: <http://www.sciencedirect.com/science/article/pii/S019688589990653X>.
- [5] R T Cox. “Probability, Frequency and Reasonable Expectation”. In: *American Journal of Physics* 14.1 (1946), pp. 1–13. DOI: 10.1119/1.1990764. URL: <https://doi.org/10.1119/1.1990764>.
- [6] B de Finetti. *Theory of Probability: A critical introductory treatment*. J. Wiley & Sons, Inc., New York, 1975.
- [7] E.T. Jaynes. “Bayesian Methods: General Background”. In: *Maximum Entropy and Bayesian Methods in Applied Statistics* August 1984 (1986), pp. 1–25. DOI: 10.1017/CB09780511569678. URL: [/chapter.jsf?bid=CB09780511569678A007%7B%5C%7Dcid=CB09780511569678A007](#).
- [8] A Stuart and K Ord. “Kendall’s Advanced Theory of Statistics: Volume I—Distribution Theory”. In: 1994.
- [9] Dirk P. Kroese and Joshua C.C. Chan. *Statistical modeling and computation*. 2014, pp. 1–400. ISBN: 9781461487753. DOI: 10.1007/978-1-4614-8775-3.
- [10] Peter McCullagh. “What is a statistical model?” In: *The Annals of Statistics* 30.5 (2002), pp. 1225–1310. ISSN: 00905364. DOI: 10.1214/aos/1035844977.
- [11] Christopher M Bishop. *Pattern Recognition and Machine Learning*. Vol. 53. 9. 2013, pp. 1689–1699. ISBN: 978-0-387-31073-2. DOI: 10.1117/1.2819119. arXiv: arXiv:1011.1669v3.
- [12] G. Upton and I. Cook. *A Dictionary of Statistics*. 2008, p. 512. ISBN: 978-0-19-954145-4. DOI: 10.1093/acref/9780199541454.001.0001.
- [13] Christopher M Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Secaucus, NJ, USA: Springer-Verlag New York, Inc., 2006. ISBN: 0387310738.
- [14] David M Blei, Alp Kucukelbir, and Jon D Mcauliffe. “Variational Inference: A Review for Statisticians”. In: (2017). arXiv: arXiv:1601.00670v6. URL: <https://arxiv.org/pdf/1601.00670.pdf>.

- [15] Giorgio Parisi. *Statistical Field Theory*. 1988.
- [16] Greg Welch and Gary Bishop. “An Introduction to the Kalman Filter”. In: (2006). URL: https://www.cs.unc.edu/%7B%7Dwelch/media/pdf/kalman%7B%5C_%7Dintro.pdf.
- [17] Shih-Tin Lin. “Force Sensing Using Kalman Filtering Techniques for Robot Compliant Motion Control”. In: *Journal of Intelligent and Robotic Systems* 18.1 (1997), pp. 1–16. ISSN: 09210296. DOI: 10.1023/A:1007946400645. URL: <http://link.springer.com/10.1023/A:1007946400645>.
- [18] David Gaylor and E. Glenn Lightsey. “GPS/INS Kalman Filter Design for Spacecraft Operating in the Proximity of International Space Station”. In: *AIAA Guidance, Navigation, and Control Conference and Exhibit*. Reston, Virgina: American Institute of Aeronautics and Astronautics, Aug. 2003. ISBN: 978-1-62410-090-1. DOI: 10.2514/6.2003-5445. URL: <http://arc.aiaa.org/doi/10.2514/6.2003-5445>.
- [19] Gurnain Kaur Pasricha. “M P RA Kalman Filter and its Economic Applications Kalman Filter and its Economic Applications”. In: (2006). URL: <http://mpira.ub.uni-muenchen.de/22734/>.
- [20] Ingvar Strid and Karl Walentin. “Block Kalman Filtering for Large-Scale DSGE Models”. In: *Computational Economics* 33.3 (Apr. 2009), pp. 277–304. ISSN: 0927-7099. DOI: 10.1007/s10614-008-9160-4. URL: <http://link.springer.com/10.1007/s10614-008-9160-4>.
- [21] S. Gannot. “Speech processing utilizing the Kalman filter”. In: *IEEE Instrumentation & Measurement Magazine* 15.3 (June 2012), pp. 10–14. ISSN: 1094-6969. DOI: 10.1109/MIM.2012.6204866. URL: <http://ieeexplore.ieee.org/document/6204866/>.
- [22] Daniel M. Wolpert and Zoubin Ghahramani. “Computational principles of movement neuroscience”. In: *Nature Neuroscience* 3.Supp (Nov. 2000), pp. 1212–1217. ISSN: 10976256. DOI: 10.1038/81497. URL: <http://www.nature.com/doifinder/10.1038/81497>.
- [23] Magdi S Mahmoud and Haris M Khalid. “Bibliographic Review on Distributed Kalman Filtering”. In: (2013). URL: <http://cogprints.org/8906/1/MsM-KFUPM-DCC-2D-One%7B%5C%7D5B3R%7B%5C%7D5D.pdf>.
- [24] Felix Govaers. “Distributed Kalman Filter”. In: *Kalman Filters - Theory for Advanced Applications*. InTech, Feb. 2018. DOI: 10.5772/intechopen.71941. URL: <http://www.intechopen.com/books/kalman-filters-theory-for-advanced-applications/distributed-kalman-filter>.
- [25] Ian F. Akyildiz and Mehmet Can Vuran. *Wireless Sensor Networks*. 2010, p. 516. ISBN: 9780470515181. DOI: 10.1002/9780470515181. URL: <http://scholar.google.com/scholar?hl=en%7B%5C%7DbtnG=Search%7B%5C%7Dq=intitle:No+Title%7B%5C%7D0%7B%5C%7D5Cnhttp://doi.wiley.com/10.1002/9780470515181>.
- [26] Kamil Dedecius and Vladimira Seckarova. “Factorized Estimation of Partially Shared Parameters in Diffusion Networks”. In: *IEEE Transactions on Signal Processing* 65.19 (Oct. 2017), pp. 5153–5163. ISSN: 1053-587X. DOI: 10.1109/TSP.2017.2725226. URL: <http://ieeexplore.ieee.org/document/7973017/>.
- [27] Mark E. Campbell and Nisar R. Ahmed. “Distributed Data Fusion: Neighbors, Rumors, and the Art of Collective Knowledge”. In: *IEEE Control Systems* 36.4 (Aug. 2016), pp. 83–109. ISSN: 1066-033X. DOI: 10.1109/MCS.2016.2558444. URL: <http://ieeexplore.ieee.org/document/7515322/>.
- [28] Ali H Sayed. “Diffusion adaptation over networks”. In: (2013). arXiv: arXiv:1205.4220v2. URL: <https://arxiv.org/pdf/1205.4220.pdf>.

- [29] Reza Olfati-Saber. “Distributed Kalman Filtering and Sensor Fusion in Sensor Networks”. In: (). URL: <https://pdfs.semanticscholar.org/fd77/043a3e588e03028589615b51f51060074a5b.pdf>.
- [30] Yulong Huang et al. “A Novel Adaptive Kalman Filter With Inaccurate Process and Measurement Noise Covariance Matrices”. In: *IEEE Transactions on Automatic Control* 63.2 (Feb. 2018), pp. 594–601. ISSN: 0018-9286. DOI: 10.1109/TAC.2017.2730480. URL: <http://ieeexplore.ieee.org/document/8025799/>.
- [31] Maurice G. (Maurice George) Kendall et al. *Kendall’s advanced theory of statistics*. Edward Arnold, 2004. ISBN: 0340807520. URL: <https://eprints.soton.ac.uk/46376/>.
- [32] Simo Särkkä Jouni Hartikainen. “Variational Bayesian Adaptation of Noise Covariances in Non-Linear Kalman Filtering”. In: (Feb. 2013). arXiv: 1302.0681. URL: <http://arxiv.org/abs/1302.0681>.